



L2b+ relative-error profiles by model

Relative error against canonical estimator (clipped at 2.5)


Opus: 88% L2b+




GPT-5: 72% L2b+




GPT-4o: 62% L2b+




Sonnet: 50% L2b+




o3: 46% L2b+



Gemini: 32% L2b+



Kimi: 10% L2b+



Reading guide

- 50% L2b+ tolerance
- 100% reference
- scenarios sorted within model

Lower, flatter profiles mean smaller numerical errors across more executed scenarios.

Scenario rank within model (sorted by relative error)