
CausalVerify: An Execution-Grounded Benchmark for LLM Causal Inference Workflows

Anonymous Author(s)

Affiliation

Address

email

Abstract

Existing causal-inference benchmarks for LLMs often evaluate text descriptions of methods or whether generated code can run. They rarely test whether the executed workflow recovers the target causal estimate. CausalVerify studies this verification problem by separating realistic interpretation from verifiable computation. The benchmark draws on 259 published economics papers, each represented by a reconstructed research question, data description, and institutional context, together with 100 fixed-seed synthetic scenarios that realise CSV datasets for difference-in-differences (DID), event study (ES), instrumental variables (IV), and regression discontinuity design (RDD). Experiment A scores method-family and direction agreement against four-LLM consensus labels. Experiment B runs model-written R code and checks whether the extracted treatment-effect estimate matches a canonical estimator on the same realised dataset. A calibration arm asks whether self-reported confidence separates correct from incorrect workflows. On Experiment B, seven evaluated LLMs reach final L2b+ pass rates between 10% and 88% at the default 50% tolerance (9% to 84% at a stricter 25%; Appendix C.4). Execution ranking (L2b) agrees with reference-agreement correctness (L2b+) far better than text-direction scoring (L4): Kendall $\tau = 0.81$ and Spearman $\rho = 0.93$, versus Kendall τ between -0.20 and 0.10 for L4. A robustness run with Llama-3.3-70B-Instruct shows the same qualitative gap, and reported confidence does not reliably separate correct from incorrect workflows. Code, data, cached outputs, and a datasheet are released with the submission.

1 Introduction

Evaluating LLMs for empirical causal inference requires more than text scoring of methods and effect signs: it requires verifying that the executed analysis recovers a target estimate. A valid empirical claim requires an estimand, an identification argument, an estimator, executable analysis code, and a numerical result whose interpretation matches the target causal quantity [52, 50, 23, 49, 27, 4]. This structure exposes a failure mode that standard LLM evaluations can miss: a model may name the right empirical design, produce code that runs without error, and still estimate the wrong coefficient. In applied research, such failures are more dangerous than obvious failures because the response appears fluent, complete, and confident while the causal claim remains unverified. This concern sits within broader verification pressure in scientific publishing, where more than 10,000 papers were retracted in 2023, a record high reported by Van Noorden [61]. As LLMs move from text assistance toward software engineering, scientific-analysis agents, and laboratory-instrumentation control [17, 66, 55, 42, 62, 24, 30, 46, 15, 63], causal-inference evaluation must ask not only whether a model sounds right, but whether its executed workflow recovers the right estimate.

Existing benchmarks capture only parts of this problem. Prior work and benchmarks on LLM causal reasoning study causal direction, graph reasoning, counterfactuals, correlation-to-causation judgments, intervention design, surveys of benchmark validity, and real-world identification tasks [34, 67, 32, 31, 14, 16, 65, 44, 68, 2, 38, 56, 54]. Code benchmarks evaluate whether programs can be synthesized, executed, or made to pass tests [13, 7, 22, 39, 43, 12, 36, 20, 30, 28, 69]. These settings are valuable, but causal workflows require a stricter check: after a model interprets a study, chooses a design, writes analysis code, and reports a treatment effect, does the executed workflow recover the intended estimate? This gap matters because code can pass surface checks while remaining semantically wrong [41].

We introduce CAUSALVERIFY, an execution-grounded benchmark for LLM causal-inference workflows. The benchmark deliberately separates ecological realism from executable verifiability. Real published papers provide realistic institutional context, but rarely provide the cleaned data, final runnable specification, and counterfactual treatment effect needed for direct verification. We therefore use 259 published economics papers to measure text-level agreement on method family and effect direction, not as verified causal correctness. Conversely, fixed-seed synthetic DGPs sacrifice some realism but provide executable targets: 100 realised CSV datasets across difference-in-differences, event study, instrumental variables, and regression discontinuity designs [25, 5, 59, 26, 37, 11, 45, 8, 1, 10, 19, 58, 53, 51]. Models write R code, the benchmark executes it, and L2b+ checks whether the extracted treatment-effect estimate matches a canonical estimator on the same realised dataset.

CAUSALVERIFY separates four signals that are often conflated: textual interpretation, code executability, numerical coefficient recovery, and self-assessed confidence. Across seven evaluated LLMs, final L2b+ pass rates range from 10% to 88%. Execution success is strongly aligned with reference-agreement correctness (Kendall $\tau = 0.81$, Spearman $\rho = 0.93$), whereas text-direction agreement is weakly or negatively aligned with Exp B L2b+ ranking (Kendall $\tau \in [-0.20, 0.10]$). Execution is not sufficient—some code runs and estimates the wrong coefficient—but it is a stronger proxy for execution-grounded agreement with the reference implementation under our reference-label construction than L4 direction agreement. A calibration arm further shows that naive retrospective confidence reports do not cleanly separate correct from incorrect workflows under our prompt, connecting the benchmark to work on calibration, verbalized uncertainty, selective prediction, and self-checking [21, 18, 29, 40, 48, 33, 35, 47, 60, 64]. Our contributions are a benchmark design separating real-paper agreement from execution-grounded reference-agreement correctness, L2b+ as a metric for agreement with a benchmark-fixed reference implementation, and empirical evidence that plausible text, runnable code, agreement with the reference implementation, and confidence can diverge sharply.

2 Related Work and Positioning

Prior work evaluates important pieces of the LLM causal-analysis workflow, but usually not the executed causal estimate. Causal benchmarks test textual reasoning about direction, counterfactuals, graphs, interventions, or identification. Code benchmarks test whether generated programs execute or pass functional tests. Science-agent benchmarks evaluate broader discovery or tool-use workflows. CAUSALVERIFY differs in its verification target: the primary correctness signal is not a text label or a generic program test, but whether executed model-written analysis code recovers a benchmark-defined target estimate on the realised dataset.

Among existing benchmarks, CauSciBench [2] is nearest in scope: it evaluates end-to-end causal-analysis performance over curated real, synthetic, and textbook tasks. CAUSALVERIFY asks a complementary measurement question: when a model appears to understand a causal task, which layer of the workflow is externally verified? It separates real-paper text agreement, controlled-DGP execution verification, and retrospective calibration rather than collapsing them into a single causal-ability score.

Concurrent work by Sawarni et al. [54] disentangles causal identification from estimation on real-world datasets, and its finding that strategy identification is easier than full specification is consistent with our name-versus-compute gap. The methodological difference is the unit of verification: CAUSALVERIFY verifies generated and executed R code on controlled DGP data, extracts the reported coefficient, and checks it against a realised-data canonical estimate. This isolates the failure




Figure 1: CAUSALVERIFY benchmark overview. Exp A uses published economics papers to form text-agreement tasks for method family and direction. Exp B uses fixed-seed synthetic DGPs and realised datasets to evaluate executable model-written R workflows against canonical estimators. The calibration arm returns each model its own Exp B answer and compares self-assessed confidence with L2b+ correctness.

mode where code runs but computes the wrong causal quantity. Recent workflow-oriented evaluations such as BioMysteryBench place models in bioinformatics environments with real data, tool access, and objective final answers [6]; CAUSALVERIFY is complementary: it fixes a realised-data causal reference estimate so that model-written causal code can be audited layer by layer rather than graded only on a final answer. Table 1 summarizes the distinction from adjacent benchmarks.

Table 1: Comparison with adjacent benchmarks. The table reports evaluation affordances rather than overall benchmark quality. A check mark (✓) indicates that a benchmark contains an explicit mechanism for that dimension, not that it is superior in all respects. ~ = partial; ✗ = absent. †EconCausal counts 10,490 triplets across 2,595 studies (Lee et al. 2025).

Benchmark	N	Target	Real ctx.	Exec. coeff.	Layered Dx	Self-assess.
CLadder / CORR2CAUSE [32, 31]	11k / 400	Causal QA / direction	~	✗	✗	✗
CausalBench [68]	3	Graph + effect	✗	✗	~	✗
EconCausal [38]	10,490†	Economic causal sign	✓	✗	✗	✗
CauSciBench [2]	367	End-to-end CI pipeline	✓	~	~	✗
InterveneBench [56]	744	Intervention design	✓	~	~	✗
CausalReasoningBench [54]	173	ID + estimate	✓	✗	~	✗
Code benchmarks [13, 30]	–	Program tests	✗	~	~	✗
Discovery / science agents [46, 15]	264 / 102	Discovery workflow	✓	~	~	✗
CAUSALVERIFY	259 + 100	Causal workflow	✓	✓	✓	✓

3 Benchmark Construction and Controls

CAUSALVERIFY separates three evaluation regimes because no single setting provides both ecological realism and executable truth. Real papers test whether models parse applied research contexts, but their ground truth is interpretive: the original cleaned data, exact runnable specification, and counterfactual effect are usually unavailable. This reflects the structure of empirical causal work, where claims depend on the alignment of estimand, identification, estimator, and data [4, 27]. Synthetic DGPs give up some realism but make coefficient correctness executable. Calibration then

102 asks the deployment-facing question: when no external verifier is available, can the model’s own
103 confidence serve as a useful warning signal?

104 **Exp A: real-paper context.** Exp A contains 259 retained published economics papers after ex-
105 cluding two source failures. Each paper is represented by three reconstructed fields: research ques-
106 tion, data description, and institutional context. Seven models are run on every paper, for 1813
107 outputs. Exp A is intentionally a text-level benchmark, comparable to prior LLM causal-reasoning
108 evaluations that score textual causal direction, method recognition, sign judgments, or identifica-
109 tion/estimation judgments [32, 31, 38, 54]. It measures method-family agreement (L3) and direc-
110 tion agreement (L4) against four-LLM consensus reference labels from a pool of Claude Opus 4.7,
111 GPT-4o, Kimi, and Gemini 2.5 Flash; this pool overlaps with the evaluated panel and we discuss the
112 implied structural circularity in §7. The benchmark does not treat these labels as human-adjudicated
113 causal claims; it uses real papers to test whether models read applied study contexts in the same
114 way as a structured consensus procedure. The execution-grounded reference-agreement endpoint is
115 reserved for Exp B.

116 **Exp B: synthetic DGP execution.** Exp B contains 100 fixed-seed synthetic DGP scenarios across
117 standard applied causal-inference designs: difference-in-differences and event studies, instrumental
118 variables, and regression discontinuity designs [10, 58, 45, 25, 5, 59, 26, 37]. The 100 scenarios
119 comprise 30 DID, 24 Event Study, 24 IV, and 22 RDD tasks; the realised datasets and canonical esti-
120 mates are fixed before model evaluation. Models receive the research question, data description, and
121 a preview of the realised CSV, then produce runnable R code. We execute the code and compare the
122 extracted treatment-effect estimate to a canonical estimator computed on the same realised dataset.
123 This realised-data comparison avoids penalizing models for finite-sample deviations between the
124 realised sample and the ideal DGP parameter. Recent large-scale multi-analyst evidence shows that
125 empirical claims can vary under alternative defensible analyses of the same research question and
126 data [57, 9, 3]. We therefore treat the Exp B canonical estimator as a fixed reference path for exe-
127 cutable evaluation, not as the only scientifically defensible analysis.

128 **Calibration: self-assessment.** We attempt a calibration query for each frozen Exp B output; re-
129 sponse availability yields 646 calibration records. The model is shown its own previous analysis
130 without the canonical estimate or pass/fail label and asked to report confidence in the method choice,
131 specification, and numerical estimate. This is a retrospective verbalized-confidence protocol rather
132 than token-probability calibration, connecting the benchmark to prior work on neural-network cali-
133 bration, LLM self-knowledge, and confidence elicitation [21, 33, 60, 64]. The resulting confidence
134 is compared to L2b+ correctness. This protocol deliberately mimics deployment: a user may see a
135 model’s confidence even when no canonical estimator is available.

136 **Controls and claim boundaries.** The benchmark uses a multi-stage evaluation pipeline rather
137 than an autonomous causal scientist. Separate stages construct contexts, audit leakage, collect
138 model responses, extract coefficients, aggregate labels, and elicit calibration, reducing “task writer
139 = respondent = grader” circularity. The synthetic scenarios are not free-form LLM inventions: math-
140 ematical templates, treatment-effect signs, sample sizes, seeds, realised CSV files, and canonical
141 estimators are fixed by the benchmark. Natural-language descriptions can vary, but a model cannot
142 pass L2b+ by guessing the sign or recognizing a phrase; it must produce code whose estimated co-
143 efficient agrees with the realised-data reference. These components support different claims: Exp A
144 is realistic but text-derived, Exp B is executable but synthetic, and calibration measures confidence-
145 correctness alignment. The benchmark is strongest when these components are read together rather
146 than collapsed into a single leaderboard.

147 4 Evaluation Layers

148 The benchmark uses six scoring labels, but they apply to different regimes. L1–L2b+ are Exp B
149 workflow-implementation layers, while L3/L4 are Exp A text-agreement diagnostics. L1 checks
150 that an output exists. L2a checks for a parseable R code block. L2b executes the code with Rscript
151 in the fixed evaluation environment. L2b+ is the execution-grounded correctness layer: it checks
152 whether the extracted treatment-effect estimate matches the realised-data canonical estimate within
153 the default relative-error tolerance. L3 measures method-family agreement and L4 measures di-

Table 2: Benchmark controls. Each component supports a different claim and protects against a different failure mode.

Component	Implementation	Claim supported
Real context	259 papers; reconstructed RQ/DD/IC; consensus reference labels	Text-level method and direction agreement
Executable truth	100 fixed-seed DGPs; realised CSVs; canonical estimators	Numerical coefficient correctness
Role separation	Separate stages for construction, response, labeling, extraction, calibration	Reduces circular grading
Leakage control	Method-name blacklist and prompt leakage audit	Prevents trivial label recovery
Scorer audit	Coefficient extraction judge plus human validation audit	Reduces brittle post-processing errors
Claim boundaries	Exp A = agreement; Exp B = execution; calibration = confidence alignment	Prevents overclaiming

rection agreement against consensus reference labels. Only L2b+ is a numerical correctness layer; L1/L2a are bookkeeping diagnostics, L2b is execution, and L3/L4 are text agreement.

Why L2b+ uses a canonical estimator. The benchmark does not ask models to recover an ideal population parameter directly. In finite samples, even a textbook estimator may differ from the DGP parameter. We therefore compare model output to a canonical estimator applied to the same realised dataset. For scalar extracted estimate $\hat{\beta}_m$ and realised-data canonical estimate $\hat{\beta}_c$, L2b+ passes when

$$\frac{|\hat{\beta}_m - \hat{\beta}_c|}{|\hat{\beta}_c|} \leq \tau,$$

with default tolerance $\tau = 0.50$. Scenarios with $|\hat{\beta}_c| < 10^{-9}$ are marked unscorable rather than divided. For Event Study scenarios, accepted post-event windows and scale conversions are frozen before scoring; valid cumulative post-event estimates are accepted only when they can be deterministically converted to the canonical per-day abnormal-return scale. This mapping is fixed before model comparison. Appendix C.1 lists the canonical reference estimate and accepted reported coefficient for each Exp B method family. L2b+ is therefore best read as agreement with a benchmark-fixed reference implementation, not as a verdict on causal-econometric correctness in general; a model that selects a methodologically more appropriate estimator (e.g. Callaway–Sant’Anna [10] or Sun–Abraham [58]) may diverge from the canonical reference, and we treat such divergence as benchmark coverage limitation rather than model error. All 30 Exp B DID scenarios use common adoption timing, in which Callaway–Sant’Anna and Sun–Abraham reduce to TWFE; Sun–Abraham implementations on three sampled scenarios pass L2b+ at 50% tolerance with relative errors below 32% (audit/rebuttal_kit/cs_sa_estimator_check.csv), confirming that the TWFE canonical reference is appropriate for this design space.

Why L3/L4 are supporting metrics. Exp A institutional context necessarily describes identifying variation, so method family can often be inferred by a trained reader. We therefore treat L3/L4 as agreement diagnostics, not verified causal correctness. They are useful diagnostics of how models parse real research contexts, but the load-bearing correctness evidence comes from Exp B.

Scorer validation. A preliminary 30-case audit of L2b-pass/L2b+-fail examples found that regex extraction produced many scorer-side failures. The final scorer therefore uses a separate coefficient-extraction judge to identify the reported treatment-effect estimate from executed stdout, followed by deterministic numerical comparison against the canonical estimate. The extraction judge sees the executed code, stdout, scenario context, and method family, but not the evaluated model identity, canonical estimate, or L2b+ label. A completed 50-cell blinded human audit of this extraction instrument finds 90.9% numeric agreement among human-scoreable effect-present cells and 88.6% agreement on the induced L2b+ pass/fail decision; Appendix E gives denominators and disagreement cases. The judge can still err, so we treat L2b+ as an audited execution-grounded metric rather than an infallible oracle.




Figure 2: Headline result. (a) L2b execution and L2b+ coefficient correctness are distinct. (b) L2b ranking tracks L2b+ correctness (Kendall $\tau = +0.81$), while L4 agreement against consensus direction labels has weak association with Exp B L2b+ ranking ($\tau \in [-0.20, +0.10]$). (c) Exp A L3/L4 agreement across the seven primary models; L3/L4 are diagnostics, not executable correctness endpoints.

188 **Threshold choice.** The default L2b+ tolerance is relative error at most 50%. This threshold is
 189 intentionally broad: it allows finite-sample and implementation variation while still rejecting esti-
 190 mates that are off by a factor large enough to change practical interpretation. The binary threshold
 191 is a coarse workflow-success endpoint; the relative-error profiles and tolerance sweep test whether
 192 conclusions depend on a single cutoff.

193 5 Results

194 The results follow the same separation as the benchmark design. Exp A first shows what real-paper
 195 text agreement can and cannot tell us. Exp B then tests the executable claim: whether the model-
 196 written workflow recovers the canonical estimate. Finally, calibration asks whether the models them-
 197 selves can identify when that workflow is wrong under a retrospective self-report prompt.

198 **Finding 1: real-paper text agreement is useful but not decisive.** On Exp A, L3 is scoreable for
 199 187/259 papers and L4 for 92/259 papers (unscorable cases are predominantly papers where the
 200 4-LLM consensus did not converge on a single canonical method-family label or positive/negative
 201 direction label). GPT-4o leads L3 (88.8%) and Sonnet leads L4 (85.9%); the rankings do not collapse
 202 to one textual ability score. The 30-paper blinded human ambiguity audit reinforces this caution:
 203 agreement with the 4-LLM consensus is only fair for direction labels (Cohen’s $\kappa = 0.294$), so L4
 204 scores partly reflect reference-label noise rather than purely model behavior. Figure 2(c) summarizes
 205 the L3/L4 separation; Figure 3 breaks L3 down by method family, and Event Study has the lowest
 206 observed cross-model L3 mean in this sample with smaller corpus size limiting precision. Exp A is
 207 therefore best read as a real-context text diagnostic, not as the paper’s primary correctness endpoint.

208 **Finding 2: runnable code is not necessarily correct.** Figure 2 shows that L2b execution and
 209 L2b+ correctness separate the model panel. Final ES-aware canonical L2b+ rates span 10–88%
 210 across the seven primary commercial models: Opus 88%, GPT-5 72%, GPT-4o 62%, Sonnet 50%,
 211 o3 46%, Gemini 32%, and Kimi 10%. This spread is invisible if evaluation stops at whether the
 212 model wrote code. Appendix C.2 decomposes Exp B failures into execution, reporting, target-
 213 coefficient, event-window, IV-stage, and RDD cutoff/sign categories.

214 **Finding 3: execution ranking tracks reference-agreement correctness.** Across the seven mod-
 215 els, the L2b ranking strongly matches the final L2b+ ranking (Kendall $\tau = +0.81$, scenario-
 216 clustered 95% bootstrap CI $[0.62, 0.90]$; Spearman $\rho = +0.93$, $[0.75, 0.96]$; 1,000 resamples, seed
 217 20260507; every replicate exceeds the L4-vs-L2b+ upper bound $+0.10$). Appendix C.3 reports
 218 leave-one-model-out stability: Kendall τ remains in $[0.733, 0.867]$, with $\tau = 0.714$ when adding
 219 Llama as robustness-only. By contrast, L4 agreement against consensus direction labels has weak
 220 or negative association with Exp B L2b+ ranking ($\tau \in [-0.20, +0.10]$ across two frozen scorer vari-
 221 ants). Execution is not sufficient, but it is a much stronger proxy for reference-agreement correctness
 222 than text-level direction. Appendix C reports $P(\text{L2b+} \mid \text{L2b})$ to separate execution failures from
 223 executed-but-wrong estimates.




Figure 3: Exp A L3 method-family agreement by method family. Colored markers show the seven model-level agreement rates; black diamonds and whiskers show the cross-model mean with paper-cluster bootstrap 95% intervals. Event Study has the lowest observed mean (64%, $n = 13$); with $n = 13$ papers the cross-model 95% bootstrap CI is wide and Event Study agreement should be treated as exploratory rather than as a comparative finding. L3 is a text-agreement diagnostic, not an executable correctness endpoint.




Figure 4: Execution cascade on Exp B. Each bar partitions 100 scenarios for one model into L2b+ correct estimates, executed-but-wrong coefficients, execution failures, and missing or unparseable code. The decomposition shows why runnable code alone is not sufficient.

Finding 4: L2b+ rankings are not threshold artifacts. The L2b+ relative-error distributions show that stronger models concentrate executed estimates near zero, while weaker models accumulate large-error tails. Thus the ranking is better understood as a distributional difference in numerical accuracy rather than an artifact of the default 50% cutoff. Appendix F reports the full relative-error profiles and tolerance sweep. Appendix C.4 reports 10%–100% tolerance sweeps; stricter cutoffs lower absolute pass rates but preserve the main text-versus-execution separation.

Finding 5: naive retrospective confidence is a weak triage signal. Calibration asks each model to rate its own previous Exp B output without seeing the L2b+ pass/fail label. Actual L2b+ rates vary sharply, but confidence gaps remain small for the main models. GPT-5 illustrates the asymmetry: it is second on L2b+ but has a slightly negative confidence gap. Under our retrospective confidence prompt, the self-assessment signal therefore does not provide a reliable deployment triage signal; this does not rule out stronger calibration interfaces. Appendix F reports reliability diagrams for the retrospective confidence protocol.

Robustness check: the execution–correctness gap reproduces for an open-weights model. As a single cross-vendor check, we evaluate Llama-3.3-70B-Instruct on the same 100 Exp B scenarios using the same prompt and scoring pipeline. It reaches L2b = 41% and final ES-aware L2b+ = 20%, placing it between Kimi and Gemini and reproducing the same qualitative gap between runnable code and numerically correct code. Adding Llama as an eighth model lowers Kendall τ to 0.714 (Appendix C.3) without changing the qualitative L2b–L2b+ ordering or L2b–L4 separation. We treat this as a robustness check, not a complete benchmark of the open-source ecosystem.

Table 3: Calibration summary against final L2b+ v2 correctness. Gap is mean numerical confidence on correct outputs minus mean confidence on incorrect outputs. Small gaps indicate weak triage value.

Model	n	L2b+ v2	Conf. gap
Opus	100	88.0	+0.075
GPT-5	99	71.7	-0.045
GPT-4o	100	62.0	+0.016
Sonnet	99	49.5	-0.011
o3	100	46.0	-0.010
Gemini	48	31.2	+0.234
Kimi	100	10.0	+0.027




Figure 5: Correctness versus self-assessment. Each point is one model. The horizontal axis is execution-grounded L2b+ correctness; the vertical axis is the mean numerical-confidence gap between correct and incorrect workflows. High L2b+ accuracy does not imply a reliable confidence signal.

244 6 Reproducibility and Human Audits

245 The final submission build is tied to a frozen audit state. The released artifacts include reconstructed
 246 Exp A paper fields, four-LLM consensus reference labels, 1813 Exp A outputs, 700 primary Exp B
 247 model-scenario execution records, 100 realised CSV datasets with canonical-estimator outputs, 646
 248 calibration records, prompt hashes, scorer decision records, and cached model responses. Determin-
 249 istic scoring can be rerun without new LLM calls; because commercial APIs are not bit-reproducible,
 250 the release treats cached JSON outputs as the reproducible model-output layer. Appendix A lists ar-
 251 tifact paths and checksums.

252 Human annotation is intentionally an ambiguity audit rather than a training or tuning input. A
 253 blinded 30-paper Exp A slice was labeled from the original PDFs after the automated benchmark
 254 freeze without access to Phase-2 LLM votes, consensus labels, or model outputs, and these human
 255 labels are not used to change frozen consensus labels, prompts, outputs, or scoring thresholds.
 256 Agreement with the four-LLM consensus is method 60.0% / direction 47.6% (Cohen’s $\kappa = 0.606$
 257 and 0.294), supporting our interpretation of Exp A as a real-context agreement diagnostic with sub-
 258 stantial ambiguity; Exp B supplies executable correctness.

259 A separate blinded human audit validates the coefficient-extraction instrument used for L2b+. The
 260 completed 50-cell audit finds 90.9% numeric agreement among human-scoreable effect-present cells
 261 and 88.6% agreement on the induced L2b+ pass/fail decision; Appendix E gives denominators and

disagreement cases. This audit validates the measurement instrument but does not change the frozen headline L2b+ rates.

7 Limitations

First, Exp A labels are reference labels rather than L2b+ correctness labels: the final labels come from a four-LLM consensus, and the completed blinded 30-paper ambiguity audit agrees with consensus on 60.0% of method labels and 47.6% of direction labels, with Cohen’s $\kappa = 0.294$ for direction. The audit is an ambiguity check, not a full human re-adjudication of the 259-paper corpus. The consensus pool also overlaps with the evaluated panel (two exact matches, two same-provider siblings; see §3), so L3/L4 are agreement diagnostics under structural overlap rather than independent text-correctness measurements; L2b+ in Exp B does not use this pool. Because Exp A uses published papers, we also cannot fully rule out model familiarity with some studies from pretraining or public metadata.

Second, Exp B is synthetic and structured. This is necessary for executable reference estimates, but it measures whether models can implement textbook workflows on known CSV data, not whether they can conduct an entire empirical research project. The DGPs cover four major method families but cannot represent all design choices in applied economics and finance. In particular, Exp B does not test literature review, data cleaning, design selection from ambiguous field constraints, or iterative robustness analysis.

Third, Exp B fixes R as the execution backend. L2b+ therefore measures executable workflow reliability under this released backend rather than language-invariant causal ability; backend sensitivity across Python, Stata, Julia, or other environments is an important extension.

Fourth, calibration is retrospective and prompt-specific. The protocol tests whether models can assess their own prior outputs under our confidence-elicitation prompt, not whether all possible uncertainty interfaces would fail. Calibration coverage is also incomplete for some models. The negative result should be read as evidence that naive retrospective self-reported confidence is not enough, not as proof that no calibrated interface can be designed.

Fifth, the evaluated model panel is limited. The primary panel consists of closed or hosted frontier systems, which is appropriate for the deployment question that motivates the paper but does not characterize the full open-source ecosystem. We add one open-weights model as a cross-vendor robustness check; a complete benchmark of multiple Llama, Qwen, Mistral, DeepSeek, and domain-specific systems is left for future work. Model-level rank correlations should also be read as descriptive diagnostics over this evaluated panel, not population-level estimates over all possible LLMs.

Sixth, L2b+ is an audited measurement instrument rather than an oracle. It depends on the canonical specification, the relative-error tolerance, event-study scale conversions, and coefficient extraction from executed stdout. We mitigate these risks through frozen scorer rules, decision records, and a completed coefficient-extraction human audit, but individual-cell mistakes remain possible and should be treated as part of the benchmark’s measurement uncertainty.

Finally, the benchmark is intentionally single-shot and non-interactive. Models receive fixed prompts and data previews, and the benchmark scores the returned workflow rather than allowing extended human-in-the-loop debugging or iterative robustness analysis. This design makes evaluation reproducible and comparable across models, but it does not measure how a supervised analyst might use an LLM over multiple rounds to improve a causal study.

8 Conclusion

CAUSALVERIFY shows that the right unit of evaluation for LLM-assisted causal inference is the workflow: a model may produce plausible causal text and runnable code while still estimating the wrong causal quantity. Across seven evaluated LLMs, final L2b+ pass rates range from 10% to 88%; execution ranking tracks reference-agreement much better than L4 direction agreement, and naive retrospective confidence does not reliably separate correct from incorrect workflows. The question for causal inference is therefore not whether the model sounds right, but whether its workflow executes, agrees with a specified target, and carries audited uncertainty.

References

- [1] Alberto Abadie. Semiparametric difference-in-differences estimators. *Review of Economic Studies*, 72(1):1–19, 2005. doi: 10.1111/0034-6527.00321.
- [2] Sawal Acharya, Terry Jingchen Zhang, Andrew Kim, Anahita Haghighat, Xianlin Sun, Rahul Babu Shrestha, Maximilian Mordig, Furkan Danisman, Clijo Jose, Yahang Qi, Pepijn Cobben, Bernhard Schölkopf, Mrinmaya Sachan, and Zhijing Jin. CauSciBench: A comprehensive benchmark on end-to-end causal inference for scientific research, 2025. URL <https://openreview.net/forum?id=uQzPkWvTy0>. Submitted to ICLR 2026.
- [3] Balazs Aczel, Barnabas Szasz, Harry T. Clelland, et al. Investigating the analytical robustness of the social and behavioural sciences. *Nature*, 652(8108):135–142, 2026. doi: 10.1038/s41586-025-09844-9. Many Analysts crowd-sourced study; 100+ contributing authors.
- [4] Joshua D. Angrist and Jörn-Steffen Pischke. *Mostly Harmless Econometrics*. Princeton University Press, 2009.
- [5] Joshua D. Angrist, Guido W. Imbens, and Donald B. Rubin. Identification of causal effects using instrumental variables. *Journal of the American Statistical Association*, 91(434):444–455, 1996. doi: 10.1080/01621459.1996.10476902.
- [6] Anthropic. Evaluating claude’s bioinformatics research capabilities with biomysterybench. <https://www.anthropic.com/research/Evaluating-Claude-For-Bioinformatics-With-BioMysteryBench>, 2026. Research blog post.
- [7] Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, and Charles Sutton. Program synthesis with large language models, 2021.
- [8] Marianne Bertrand, Esther Duflo, and Sendhil Mullainathan. How much should we trust differences-in-differences estimates? *Quarterly Journal of Economics*, 119(1):249–275, 2004. doi: 10.1162/003355304772839588.
- [9] R. Botvinik-Nezer, F. Holzmeister, C. F. Camerer, et al. Variability in the analysis of a single neuroimaging dataset by many teams. *Nature*, 582(7810):84–88, 2020. doi: 10.1038/s41586-020-2314-9.
- [10] Brantly Callaway and Pedro H. C. Sant’Anna. Difference-in-differences with multiple time periods. *Journal of Econometrics*, 225(2):200–230, 2021. doi: 10.1016/j.jeconom.2020.12.001.
- [11] Sebastian Calonico, Matias D. Cattaneo, and Rocio Titiunik. Robust nonparametric confidence intervals for regression-discontinuity designs. *Econometrica*, 82(6):2295–2326, 2014. doi: 10.3982/ECTA11757.
- [12] Federico Cassano, John Gouwar, Daniel Nguyen, Sydney Nguyen, Luna Phipps-Costin, Donald Pinckney, Ming-Ho Yee, Yangtian Zi, Carolyn Jane Anderson, Molly Q. Feldman, Arjun Guha, Michael Greenberg, and Abhinav Jangda. MultiPL-E: A scalable and polyglot approach to benchmarking neural code generation. *IEEE Transactions on Software Engineering*, 2023. doi: 10.1109/TSE.2023.3267446.
- [13] Mark Chen, Jerry Tworek, et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- [14] Sirui Chen, Mengying Xu, Kun Wang, Xingyu Zeng, Rui Zhao, Shengjie Zhao, and Chaochao Lu. CLEAR: Can language models really understand causal graphs? In *Findings of the Association for Computational Linguistics: EMNLP 2024*, 2024. doi: 10.18653/v1/2024.findings-emnlp.363.

- [15] Zirui Chen, Shijie Chen, Yuting Ning, Qianheng Zhang, Boshu Wang, Botao Yu, Yifei Li, Zeyi Liao, Chen Wei, Zitong Lu, Vishal Dey, Frazier N. Baker, Benjamin Burns, Daniel Adu-Ampratwum, Xuhui Huang, Xia Ning, Song Gao, Yu Su, and Huan Sun. ScienceAgentBench: Toward rigorous assessment of language agents for data-driven scientific discovery. In *International Conference on Learning Representations (ICLR)*, 2025.
- [16] Haoang Chi, He Li, Wenjing Yang, Feng Liu, Long Lan, Xiaoguang Ren, Tongliang Liu, and Bo Han. Unveiling causal reasoning in large language models: Reality or mirage? In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi: 10.52202/079017-3064.
- [17] Luyu Gao, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, and Graham Neubig. PAL: Program-aided language models. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 10764–10799. PMLR, 2023. doi: 10.48550/arXiv.2211.10435.
- [18] Yonatan Geifman and Ran El-Yaniv. Selective classification for deep neural networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [19] Andrew Goodman-Bacon. Difference-in-differences with variation in treatment timing. *Journal of Econometrics*, 225(2):254–277, 2021. doi: 10.1016/j.jeconom.2021.03.014.
- [20] Alex Gu, Baptiste Rozière, Hugh Leather, Armando Solar-Lezama, Gabriel Synnaeve, and Sida Wang. CRUXEval: A benchmark for code reasoning, understanding and execution. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *Proceedings of Machine Learning Research*, pages 16568–16621. PMLR, 2024. URL <https://proceedings.mlr.press/v235/gu24c.html>.
- [21] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *International Conference on Machine Learning (ICML)*, 2017.
- [22] Dan Hendrycks, Steven Basart, Saurav Kadavath, Mantas Mazeika, Akul Arora, Ethan Guo, Collin Burns, Samir Puranik, Horace He, Dawn Song, and Jacob Steinhardt. Measuring coding challenge competence with APPS. In *Advances in Neural Information Processing Systems*, volume 34, 2021. doi: 10.48550/arXiv.2105.09938.
- [23] Paul W. Holland. Statistics and causal inference. *Journal of the American Statistical Association*, 81(396):945–960, 1986. doi: 10.1080/01621459.1986.10478354.
- [24] Qian Huang, Jian Vora, Percy Liang, and Jure Leskovec. MLAGentBench: Evaluating language agents on machine learning experimentation. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, volume 235 of *Proceedings of Machine Learning Research*, pages 20271–20309. PMLR, 2024.
- [25] Guido W. Imbens and Joshua D. Angrist. Identification and estimation of local average treatment effects. *Econometrica*, 62(2):467–475, 1994. doi: 10.2307/2951620.
- [26] Guido W. Imbens and Thomas Lemieux. Regression discontinuity designs: A guide to practice. *Journal of Econometrics*, 142(2):615–635, 2008. doi: 10.1016/j.jeconom.2007.05.001.
- [27] Guido W. Imbens and Donald B. Rubin. *Causal Inference for Statistics, Social, and Biomedical Sciences: An Introduction*. Cambridge University Press, Cambridge, UK, 2015. doi: 10.1017/CBO9781139025751.
- [28] Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. LiveCodeBench: Holistic and contamination free evaluation of large language models for code. In *International Conference on Learning Representations (ICLR)*, 2025. URL <https://openreview.net/forum?id=chfJJYC3iL>.
- [29] Zhengbao Jiang, Jun Araki, Haibo Ding, and Graham Neubig. How can we know when language models know? on the calibration of language models for question answering. *Transactions of the Association for Computational Linguistics*, 9:962–977, 2021. doi: 10.1162/tacl_a_00407.

- [30] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. Swe-bench: Can language models resolve real-world GitHub issues? In *ICLR*, 2024.
- [31] Zhijing Jin, Jiarui Liu, Zhiheng Lyu, Spencer Poff, Mrinmaya Sachan, Rada Mihalcea, Mona Diab, and Bernhard Schölkopf. Can large language models infer causation from correlation? In *International Conference on Learning Representations*, 2024. doi: 10.48550/arXiv.2306.05836. URL <https://openreview.net/forum?id=vqIH00bdqL>.
- [32] Zhijing Jin et al. CLadder: Assessing causal reasoning in language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023. doi: 10.52202/075280-1353.
- [33] Saurav Kadavath, Tom Conerly, Amanda Askell, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- [34] Emre Kiciman, Robert Ness, Amit Sharma, and Chenhao Tan. Causal reasoning and large language models: Opening a new frontier for causality. *Transactions on Machine Learning Research*, 2024. ISSN 2835-8856.
- [35] Lorenz Kuhn, Yarin Gal, and Sebastian Farquhar. Semantic uncertainty: Linguistic invariances for uncertainty estimation in natural language generation. In *International Conference on Learning Representations (ICLR)*, 2023. Oral.
- [36] Yuhang Lai, Chengxi Li, Yiming Wang, Tianyi Zhang, Ruiqi Zhong, Luke Zettlemoyer, Scott Wen-tau Yih, Daniel Fried, Sida Wang, and Tao Yu. DS-1000: A natural and reliable benchmark for data science code generation. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 18319–18345. PMLR, 2023. doi: 10.48550/arXiv.2211.11501.
- [37] David S. Lee and Thomas Lemieux. Regression discontinuity designs in economics. *Journal of Economic Literature*, 48(2):281–355, 2010. doi: 10.1257/jel.48.2.281.
- [38] Donggyu Lee, Hyeok Yun, Meeyoung Cha, Sungwon Park, Sangyoon Park, and Jihee Kim. EconCausal: A context-aware causal reasoning benchmark for large language models in social science, 2025.
- [39] Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, Thomas Hubert, Peter Choy, Cyprien de Masson d’Autume, Igor Babuschkin, Xinyun Chen, Po-Sen Huang, Johannes Welbl, Sven Gowal, Alexey Cherepanov, James Molloy, Daniel J. Mankowitz, Esme Sutherland Robson, Pushmeet Kohli, Nando de Freitas, Koray Kavukcuoglu, and Oriol Vinyals. Competition-level code generation with AlphaCode. *Science*, 378(6624):1092–1097, 2022. doi: 10.1126/science.abq1158.
- [40] Stephanie Lin, Jacob Hilton, and Owain Evans. Teaching models to express their uncertainty in words. *Transactions on Machine Learning Research*, 2022. doi: 10.48550/arXiv.2205.14334. URL <https://openreview.net/forum?id=8s8K2UZGTZ>.
- [41] Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and Lingming Zhang. Is your code generated by ChatGPT really correct? rigorous evaluation of large language models for code generation. In *Advances in Neural Information Processing Systems*, volume 36, 2023. doi: 10.48550/arXiv.2305.01210.
- [42] Xiao Liu, Hao Yu, Hanchen Zhang, Yifan Xu, Xuanyu Lei, Hanyu Lai, Yu Gu, Hangliang Ding, Kaiwen Men, Kejuan Yang, Shudan Zhang, Xiang Deng, Aohan Zeng, Zhengxiao Du, Chenhui Zhang, Sheng Shen, Tianjun Zhang, Yu Su, Huan Sun, Minlie Huang, Yuxiao Dong, and Jie Tang. AgentBench: Evaluating LLMs as agents. In *International Conference on Learning Representations (ICLR)*, 2024.
- [43] Shuai Lu, Daya Guo, Shuo Ren, Junjie Huang, Alexey Svyatkovskiy, Ambrosio Blanco, Colin B. Clement, Dawn Drain, Daxin Jiang, Duyu Tang, Ge Li, Lidong Zhou, Linjun Shou, Long Zhou, Michele Tufano, Ming Gong, Ming Zhou, Nan Duan, Neel Sundaresan, Shaowei Deng, Shengyu Fu, and Shujie Liu. CodeXGLUE: A machine learning benchmark dataset for

- code understanding and generation. In *Advances in Neural Information Processing Systems Track on Datasets and Benchmarks*, 2021.
- [44] Jing Ma. Causal inference with large language model: A survey. In *Findings of the Association for Computational Linguistics: NAACL 2025*, 2025. doi: 10.18653/v1/2025.findings-naacl.327.
- [45] A. Craig MacKinlay. Event studies in economics and finance. *Journal of Economic Literature*, 35(1):13–39, 1997.
- [46] Bodhisattwa Prasad Majumder, Harshit Surana, Dhruv Agarwal, Bhavana Dalvi Mishra, Abhisheksingh Meena, Aryan Prakhar, Tirth Vora, Tushar Khot, Ashish Sabharwal, and Peter Clark. DiscoveryBench: Towards data-driven discovery with large language models. In *International Conference on Learning Representations (ICLR)*, 2025.
- [47] Potsawee Manakul, Adian Liusie, and Mark J. F. Gales. SelfCheckGPT: Zero-resource black-box hallucination detection for generative large language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 9004–9017, 2023. doi: 10.18653/v1/2023.emnlp-main.557.
- [48] Sabrina J. Mielke, Arthur Szlam, Emily Dinan, and Y-Lan Boureau. Reducing conversational agents’ overconfidence through linguistic calibration. *Transactions of the Association for Computational Linguistics*, 10:857–872, 2022. doi: 10.1162/tacl_a_00494.
- [49] Judea Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, Cambridge, UK, 2 edition, 2009. doi: 10.1017/CBO9780511803161.
- [50] Paul R. Rosenbaum and Donald B. Rubin. The central role of the propensity score in observational studies for causal effects. *Biometrika*, 70(1):41–55, 1983. doi: 10.1093/biomet/70.1.41.
- [51] Jonathan Roth, Pedro H. C. Sant’Anna, Alyssa Bilinski, and John Poe. What’s trending in difference-in-differences? a synthesis of the recent econometrics literature. *Journal of Econometrics*, 235(2):2218–2244, 2023. doi: 10.1016/j.jeconom.2023.03.008.
- [52] Donald B. Rubin. Estimating causal effects of treatments in randomized and nonrandomized studies. *Journal of Educational Psychology*, 66(5):688–701, 1974. doi: 10.1037/h0037350.
- [53] Pedro H. C. Sant’Anna and Jun B. Zhao. Doubly robust difference-in-differences estimators. *Journal of Econometrics*, 219(1):101–122, 2020. doi: 10.1016/j.jeconom.2020.06.003.
- [54] Ayush Sawarni, Jiyuan Tan, and Vasilis Syrgkanis. CausalReasoningBenchmark: A real-world benchmark for disentangled evaluation of causal identification and estimation, 2026.
- [55] Timo Schick, Jane Dwivedi-Yu, Roberto Dessi, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems*, volume 36, 2023. doi: 10.48550/arXiv.2302.04761.
- [56] Shaojie Shi, Zhengyu Shi, Lingran Zheng, Xinyu Su, Anna Xie, Bohao Lv, Rui Xu, Zijian Chen, Zhichao Chen, Guolei Liu, Naifu Zhang, Mingjian Dong, Zhuo Quan, Bohao Chen, Teqi Hao, Yuan Qi, Yinghui Xu, and Libo Wu. InterveneBench: Benchmarking LLMs for intervention reasoning and causal study design in real social systems, 2026.
- [57] R. Silberzahn, E. L. Uhlmann, D. P. Martin, et al. Many analysts, one data set: Making transparent how variations in analytic choices affect results. *Advances in Methods and Practices in Psychological Science*, 1(3):337–356, 2018. doi: 10.1177/2515245917747646.
- [58] Liyang Sun and Sarah Abraham. Estimating dynamic treatment effects in event studies with heterogeneous treatment effects. *Journal of Econometrics*, 225(2):175–199, 2021. doi: 10.1016/j.jeconom.2020.09.006.
- [59] Donald L. Thistlethwaite and Donald T. Campbell. Regression-discontinuity analysis: An alternative to the ex post facto experiment. *Journal of Educational Psychology*, 51(6):309–317, 1960. doi: 10.1037/h0044319.

- [60] Katherine Tian, Eric Mitchell, Allan Zhou, Archit Sharma, Rafael Rafailov, Huaxiu Yao, Chelsea Finn, and Christopher D. Manning. Just ask for calibration: Strategies for eliciting calibrated confidence scores from language models fine-tuned with human feedback. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 5433–5442, 2023. doi: 10.18653/v1/2023.emnlp-main.330.
- [61] Richard Van Noorden. More than 10,000 research papers were retracted in 2023—a new record. *Nature*, 624(7992):479–481, 2023. doi: 10.1038/d41586-023-03974-8. URL <https://pubmed.ncbi.nlm.nih.gov/38087103/>.
- [62] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, Ahmed Hassan Awadallah, Ryen W. White, Doug Burger, and Chi Wang. AutoGen: Enabling next-gen LLM applications via multi-agent conversation, 2023.
- [63] Yong Xie, Kexin He, and Andres Castellanos-Gomez. Toward full autonomous laboratory instrumentation control with large language models. *Small Structures*, 6:2500173, 2025. doi: 10.1002/sstr.202500173.
- [64] Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs. In *International Conference on Learning Representations*, 2024. doi: 10.48550/arXiv.2306.13063. URL <https://openreview.net/forum?id=gjeQKFxFpZ>.
- [65] Linying Yang, Vik Shirvaikar, Oscar Clivio, and Fabian Falck. A critical review of causal reasoning benchmarks for large language models, 2024.
- [66] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations*, 2023. doi: 10.48550/arXiv.2210.03629.
- [67] Matej Zecevic, Moritz Willig, Devendra Singh Dhami, and Kristian Kersting. Causal parrots: Large language models may talk causality but are not causal. *Transactions on Machine Learning Research*, 2023. URL <https://openreview.net/forum?id=tv46tCzs83>.
- [68] Yu Zhou, Xingyu Wu, Beicheng Huang, Jibin Wu, Liang Feng, and Kay Chen Tan. Causal-Bench: A comprehensive benchmark for causal learning capability of large language models, 2024.
- [69] Terry Yue Zhuo, Minh Chien Vu, Jenny Chim, Han Hu, Wenhao Yu, et al. BigCodeBench: Benchmarking code generation with diverse function calls and complex instructions. In *International Conference on Learning Representations (ICLR)*, 2025. doi: 10.48550/arXiv.2406.15877. Oral.

A Reproducibility Details

Appendix roadmap. Appendix A lists frozen artifacts and checksums. Appendix B describes the 30-paper Exp A ambiguity audit. Appendix C reports Exp B robustness diagnostics: canonical-estimator targets, conditional $P(\text{L2b+} \mid \text{L2b})$, rank-stability, tolerance sweep, scorer evolution, and method-family breakdown. Appendix D contains the descriptive RID pilot. Appendix E reports the 50-cell coefficient-extraction human audit. Appendix F contains additional diagnostic figures moved from the main text.

The final submission build is tied to a complete audit state. The audit tag `v11-freeze-2026-04-29` preserves the full provenance trail, while `causalverify_neurips2026.tex` provides a submission-length rendering of the same claims. The final submission PDF checksum is documented in `audit/SUBMISSION_BUILD_SUMMARY.md`.

Table 4: Frozen artifacts used by the final submission build. The full audit state remains available for inspection; the submission paper is a shorter rendering of the same frozen claims.

Claim	Frozen artifact
Exp A outputs	<code>experiments/exp_a/outputs/</code>
Exp A scores	<code>experiments/exp_a/auto_scores.csv</code>
Exp B L2b+	<code>l2b_plus_scores_canonical_judge_v2.csv</code> (primary 7; Llama R6 retained)
Calibration	<code>calibration_scores.csv</code> , <code>calibration_summary_v2.json</code>
Human validation	<code>audit/human_gold/human_vs_llm_consensus.md</code>
Paper checksum	<code>audit/SUBMISSION_BUILD_SUMMARY.md</code>
Gate evidence	<code>audit/V11_ACCEPTANCE_GATES.md</code>

B Human Annotation Details

The repository contains a paper-native labeling protocol and a completed blinded 30-paper ambiguity audit slice. The annotator labeled method family and direction from the original PDFs without access to Phase-2 LLM votes, consensus labels, or model outputs. One target paper (`paper_159`) was excluded because the available PDF was appendix-only; `paper_17` was added as a blinded replacement from the unsealed sample manifest, which contains only paper ID, difficulty tier, and PDF path. These human labels are not used to silently change the frozen Exp A labels. Human adjudication was conducted after the automated benchmark freeze as an ambiguity audit, not as a training, prompting, or tuning input; the labels are not used to construct prompts, select model outputs, or tune scoring thresholds.

C Exp B Robustness Details

This appendix reports supporting diagnostics for Exp B. These tables are derived from `l2b_plus_scores_canonical_judge_v2.csv` and do not replace the frozen headline L2b+ rates in the main text.

C.1 Canonical estimators for Exp B

L2b+ compares the model-reported estimate to a fixed realised-data reference estimate. The canonical estimator is a benchmark-defined reference path for executable evaluation, not a claim that only one empirical analysis is scientifically defensible. It is also distinct from the structural DGP parameter used to simulate the data.

Because L2b+ requires code execution, raw L2b–L2b+ association is partly structural. Table 6 therefore reports the conditional metric $P(\text{L2b+} \mid \text{L2b})$, which asks whether executed code computes the right coefficient.

Table 5: Benchmark-defined canonical estimators for Exp B L2b+ scoring.

Method	Canonical reference estimate	Accepted coefficient	model-reported	Common non-target coefficients	Scope limitation
DID	TWFE coefficient on <code>treat_x_post</code> after unit and period demeaning	on <code>treat_x_post</code> , <code>treated:post</code> , or equivalent treatment-by-post effect	post-event	<code>treated</code> main effect, post main effect, placebo or pre-trend coefficient	Clean benchmark DID template, not a full benchmark of staggered or heterogeneous-treatment DID estimators
Event Study	Market-model coefficient on <code>event_day >= 0</code> ; ES-aware v2 accepts frozen integer post-window conversions to this per-day scale	Explicit AAR/CAR that can be deterministically converted to the canonical scale	post-event	Pre-event coefficient, arbitrary non-frozen window, diagnostic or placebo return	Fixed benchmark window mapping, not a universal event-study estimand
IV	Realised-data 2SLS coefficient on <code>x</code> , instrumented by <code>z</code> , with controls exogenous when present	Second-stage coefficient on <code>x</code> or the named endogenous variable		First-stage coefficient on <code>z</code> , reduced form, or control coefficient	Tests implementation of the IV estimand; does not validate exclusion restriction or weak-instrument diagnostics
RDD	Local-linear cutoff effect on the realised data: coefficient on <code>treated</code> near cutoff 0.5 with running-variable slope and interaction	Treatment, above-cutoff, or discontinuity coefficient consistent with benchmark direction		Density test, bandwidth statistic, unrelated global polynomial trend, placebo cutoff estimate	Fixed benchmark RDD reference, not an exhaustive comparison of bandwidth or inference choices

Table 6: Conditional Exp B correctness among executed outputs. “Wrong after execution” is L2b count minus L2b+ v2 count.

Model	L2b	L2b+	Wrong after exec.	$P(\text{L2b+} \text{L2b})$	Gap
Opus	94	88	6	93.6%	6 pp
GPT-5	76	72	4	94.7%	4 pp
GPT-4o	78	62	16	79.5%	16 pp
Sonnet	51	50	1	98.0%	1 pp
o3	50	46	4	92.0%	4 pp
Gemini	32	32	0	100.0%	0 pp
Kimi	45	10	35	22.2%	35 pp

572 C.2 Exp B failure-taxonomy diagnostic

573 We also audit the non-L2b+ mass with a conservative rule-based taxonomy. The taxonomy is diag-
574 nostic rather than a scoring replacement: it distinguishes missing or unparseable code, execution fail-
575 ures, executed-but-unreported effects, wrong target coefficients, event-window or scale mismatches,
576 IV stage errors, and RDD cutoff or sign-convention issues. Ambiguous cases are left uncategorized
577 rather than forced into a narrow label.

Table 7: Exp B failure-taxonomy diagnostic, per primary model. Cells show counts and (share of that model’s non-L2b+ failures). Per-model breakdown reveals that strong models execute correctly when they execute (Opus, Sonnet, Gemini conditional pass rates $\geq 93\%$; see Table 6), while weaker models more often produce executed-but-wrong code (Kimi 22.2%, GPT-4o 79.5% conditional pass). Per-model rows sum to 340 non-L2b+ cells, matching the aggregate decomposition (Exec fail 230, No code 44, Wrong coef 26, Coef-not-reported 15, Window/scale 11, IV stage 5, RDD 3, Other 6).

Model	n	Exec fail	No code	Wrong coef	Coef N/R	Win/scale	IV stage	RDD	Other
Opus	12	5 (41.7%)	1 (8.3%)	2 (16.7%)	0	3 (25.0%)	0	0	1 (8.3%)
Sonnet	50	49 (98.0%)	0	1 (2.0%)	0	0	0	0	0
GPT-4o	38	22 (57.9%)	0	3 (7.9%)	1 (2.6%)	3 (7.9%)	4 (10.5%)	0	5 (13.2%)
o3	54	50 (92.6%)	0	2 (3.7%)	1 (1.9%)	1 (1.9%)	0	0	0
Kimi	90	55 (61.1%)	0	18 (20.0%)	13 (14.4%)	3 (3.3%)	1 (1.1%)	0	0
Gemini	68	25 (36.8%)	43 (63.2%)	0	0	0	0	0	0
GPT-5	28	24 (85.7%)	0	0	0	1 (3.6%)	0	3 (10.7%)	0

578 C.3 Rank-stability diagnostic

579 The rank correlations in Figure 2 are descriptive model-level diagnostics over seven primary models
 580 rather than population-level estimates. To assess panel sensitivity, we recompute Kendall τ and
 581 Spearman ρ after leaving out each model in turn. The leave-one-model-out Kendall range is 0.733–
 582 0.867 and the Spearman range is 0.886–0.943. Adding Llama-3.3-70B-Instruct as a robustness-only
 583 eighth model gives Kendall $\tau = 0.714$ and Spearman $\rho = 0.881$, while the primary leaderboard
 584 remains the seven-model panel.

Table 8: Rank-stability diagnostic for the Exp B model-level correlations. Primary-panel rows summarize the audit table in `paper/tables/exp_b_rank_stability_primary7.csv`; the Llama row is robustness-only and does not enter the primary leaderboard.

Diagnostic	N models	Kendall τ	Spearman ρ
Primary seven-model baseline	7	0.810	0.929
Leave-one-model-out range	6 each	0.733–0.867	0.886–0.943
Primary seven plus Llama robustness-only	8	0.714	0.881

Table 9: Scorer evolution on Exp B. Regex extraction is retained as a legacy diagnostic; the final headline uses the ES-aware v2 score. Entries are pass rates in percentage points.

Model	Regex	Judge	ES-aware v2	v2–Judge
Opus	64	71	88	+17
GPT-5	29	59	72	+13
GPT-4o	28	50	62	+12
Sonnet	27	31	50	+19
o3	18	37	46	+9
Gemini	15	29	32	+3
Kimi	6	8	10	+2

Table 10: L2b+ v2 by model and method family. Each cell reports passes/available scenarios.

Model	DID	ES	IV	RDD
Opus	29/30=97%	19/24=79%	23/24=96%	17/22=77%
GPT-5	20/30=67%	17/24=71%	23/24=96%	12/22=55%
GPT-4o	25/30=83%	13/24=54%	15/24=62%	9/22=41%
Sonnet	20/30=67%	21/24=88%	3/24=12%	6/22=27%
o3	8/30=27%	12/24=50%	16/24=67%	10/22=45%
Gemini	14/30=47%	5/24=21%	13/24=54%	0/22=0%
Kimi	5/30=17%	4/24=17%	1/24=4%	0/22=0%

585 C.4 Tolerance-sweep diagnostic

586 The binary L2b+ endpoint uses a broad default relative-error tolerance of 50%. To check whether the
 587 qualitative conclusion depends on this cutoff, we recompute primary-seven pass rates at 10%, 25%,
 588 50%, 75%, and 100%. The sweep is used as a robustness diagnostic; the frozen headline results
 589 continue to use the default 50% tolerance. Table 11 reports all-scenario pass rates, so non-executing
 590 outputs remain failures.

Table 11: Tolerance-sweep diagnostic for Exp B L2b+ scoring. Each cell reports the pass rate out of 100 scenarios; parentheses give the model rank at that tolerance. The 50% column is the frozen headline threshold.

Model	10%	25%	50%	75%	100%
Opus	72% (1)	84% (1)	88% (1)	88% (1)	88% (1)
GPT-5	64% (2)	71% (2)	72% (2)	72% (2)	72% (2)
GPT-4o	53% (3)	59% (3)	62% (3)	63% (3)	63% (3)
Sonnet	42% (4)	48% (4)	50% (4)	50% (4)	50% (4)
o3	38% (5)	44% (5)	46% (5)	47% (5)	47% (5)
Gemini	32% (6)	32% (6)	32% (6)	32% (6)	32% (6)
Kimi	9% (7)	9% (7)	10% (7)	13% (7)	14% (7)

591 D RID Pilot Ablation

592 Experiment C is a small descriptive pilot that reruns a 10-paper subset under the RID pre-
 593 commitment protocol. It is not used for the headline benchmark claims, because $N = 10$ per
 594 model is too small for model-level inference. We include it only as an audit artifact showing how
 595 pre-commitment can change text-level L3/L4 outcomes in either direction.




Figure 6: RID pre-commitment ablation pilot ($N = 10$ papers per model, descriptive only). Dots compare RID OFF and RID ON pass rates for L3 strategy and L4 direction. No individual comparison is statistically significant; the pilot is retained as an audit artifact rather than as evidence for the main claims.

596 E Coefficient-Extraction Human Audit

597 The coefficient-extraction judge is a measurement instrument: it identifies the scalar treatment-
 598 effect coefficient from executed R code/stdout before L2b+ compares that coefficient to
 599 the canonical estimator. The repository contains a blinded human-validation audit in
 600 audit/l2b_judge_human_validation/. The annotation form hides model identity, L2b+ label,
 601 canonical estimate, judge effect, regex effect, and ranking information. The completed audit covers
 602 50 primary-panel L2b=1 cells sampled with seed 20260502: DID 13, Event Study 12, IV 11, and
 603 RDD 14. Among comparable audited cells, it found 90.9% numeric agreement and 88.6% agree-
 604 ment on the induced L2b+ pass/fail decision. Disagreements concentrate in event-study window
 605 choices and RDD sign/printing ambiguities. This audit validates the coefficient-extraction instru-
 606 ment; it does not change the frozen headline L2b+ rates.

607 **Disagreement categories.** Most disagreements arise from two sources: (i) Event Study outputs
 608 that report a cumulative post-event window rather than the canonical per-day scale, and (ii) RDD
 609 outputs where printed signs, jump orientation, or robustness estimates make the main effect am-
 610 biguous. These cases motivate treating L2b+ as an audited measurement instrument rather than an
 611 infallible oracle.

612 F Additional Diagnostics

613 This appendix collects supporting diagnostic figures for the main results. The main text reports
 614 the headline conclusions; the relative-error profiles, calibration reliability diagrams, and L4 scorer-
 615 operationalization diagnostic below support those conclusions but are not required for the central
 616 correctness argument.




Figure 7: Additional Exp B robustness diagnostic. Each panel sorts executed estimates by relative error against the realised-data canonical estimator; the dashed line marks the default 50% tolerance. These profiles support the main-text claim that L2b+ rankings are not driven by a single cutoff.




Figure 8: Additional calibration diagnostic. Reliability diagrams compare retrospective self-reported confidence with empirical L2b+ correctness. These diagnostics support the main-text finding that naive retrospective self-reported confidence does not reliably separate correct from incorrect workflows under our retrospective confidence prompt. 95% bootstrap CI on ECE (1,000 resamples within model): see `paper/derived_analyses/ece_bootstrap_ci.csv`.




Figure 9: L4 scorer-operationalization diagnostic. Alternative frozen L4 text-direction scorers produce unstable pass rates across models. This diagnostic is not part of the primary seven-model headline figure because the scorer sweep was run on the six models with frozen scorer-variant outputs.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction define the scope as a benchmark for verifying LLM causal-inference workflows. They distinguish real-paper text agreement from execution-grounded coefficient correctness, and state that Exp A measures agreement diagnostics while Exp B provides the executable correctness endpoint.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: The Limitations section covers Exp A label ambiguity, the 30-paper human ambiguity audit, the fact that Exp A is not a fully human-adjudicated reference-label corpus, the synthetic scope of Exp B, the dependence of L2b+ on canonical estimators and coefficient extraction, model snapshot limits, and domain scope.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (correct) proof?

Answer: [N/A]

Justification: This is an empirical benchmark paper; no theoretical results or proofs are presented.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper?

Answer: [Yes]

Justification: The anonymized release includes 259 active Exp A paper records, four-LLM consensus labels, 100 Exp B synthetic DGP scenarios with realised CSV data files and canonical-estimator targets, 1813 Exp A outputs, 700 Exp B primary execution cells, 646 calibration records, the L2b+ coefficient-extraction pipeline, prompt templates, scoring scripts, and frozen audit summaries.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient documentation and instructions for reproducing the main experimental results?

Answer: [Yes]

Justification: The anonymized repository submitted with the paper contains the code, frozen outputs, documentation, datasheet, component-level data-license notes, and reproduction commands. The README explains how to rerun deterministic scoring for Exp A, Exp B, and calibration without making new LLM API calls.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: [Yes]

Justification: No training is performed; this is an evaluation-only benchmark. The paper and release specify the evaluated models, prompt construction, Exp A consensus-label construction, Exp B DGP generation, realised-data canonical estimators, L2b/L2b+ scoring, and calibration protocol. Prompt templates and implementation details are included in the code release.

667 **7. Experiment statistical significance**

668 Question: Does the paper report error bars suitably and correctly defined or other appropri-
669 ate information about the statistical significance of the experiments?

670 Answer: [Yes]

671 Justification: The paper reports sample sizes for all scoreable subsets, model-level pass
672 rates, rank correlations for the L2b/L2b+ and L4/L2b+ comparisons, and a relative-error
673 distribution analysis showing that the main Exp B ranking is not an artifact of a single
674 tolerance threshold. The paper does not make pairwise model-significance claims.

675 **8. Experiments compute resources**

676 Question: For each experiment, does the paper provide sufficient information on the com-
677 puter resources (type, memory, time, etc.) needed to reproduce the experiments?

678 Answer: [Yes]

679 Justification: The original model-generation runs use commercial LLM API calls and no
680 local GPU compute. The deterministic scoring and figure-generation steps can be rerun
681 from frozen outputs on a standard CPU environment. Token counts, costs, prompts, and
682 run logs are preserved in the audit artifacts.

683 **9. Code of ethics**

684 Question: Does the research conducted in the paper conform, in every respect, with the
685 NeurIPS Code of Ethics?

686 Answer: [Yes]

687 Justification: The paper evaluates publicly available LLMs on published economics papers
688 and synthetic scenarios. No human subjects are involved. The broader impact paragraph
689 (Section 8) discusses risks of using LLMs for policy-facing causal analysis.

690 **10. Broader impacts**

691 Question: Does the paper discuss both potential positive societal impacts and negative
692 societal impacts of the work?

693 Answer: [Yes]

694 Justification: Section 8 (Broader impact paragraph) discusses the risk that LLMs used in
695 policy-facing causal analysis without recognizing documented failure modes could support
696 misleading conclusions. The benchmark itself is designed to surface these failure modes,
697 which is a positive impact.

698 **11. Safeguards**

699 Question: Does the paper describe safeguards that have been put in place for responsible
700 release of data or models that have a high risk for misuse?

701 Answer: [N/A]

702 Justification: The released data consists of published paper metadata and synthetic scenar-
703 ios; the released code is an evaluation framework. Neither poses a meaningful misuse risk.

704 **12. Licenses for existing assets**

705 Question: Are the creators or original owners of assets (e.g., code, data, models) used in
706 the paper properly credited and are the license and terms of use explicitly mentioned and
707 properly respected?

708 Answer: [Yes]

709 Justification: The 259 active paper records include source metadata and provenance. Mod-
710 els are identified by provider and API version in the release documentation. The code is
711 released under the repository's code license, synthetic DGP artifacts and scored outputs
712 are licensed as documented in the data-license file, and original published article PDFs
713 are not relicensed by this repository unless redistribution rights are explicitly confirmed.
714 OpenAlex-derived metadata is used under its stated terms.

715 **13. New assets**

716 Question: Are new assets introduced in the paper well documented and is the documenta-
717 tion provided alongside the assets in a format that is easily accessible?

718 Answer: [\[Yes\]](#)

719 Justification: The benchmark is accompanied by a DATASHEET.md following Gebru et al.
 720 (2021), a detailed README with reproduction instructions, and structured JSON schemas
 721 for all papers and scenarios.

722 **14. Crowdsourcing and research with human subjects**

723 Question: For crowdsourcing experiments and research with human subjects, does the pa-
 724 per include the full text of instructions given to participants and screenshots?

725 Answer: [\[N/A\]](#)

726 Justification: No crowdsourcing or human subjects research was conducted. A 30-paper
 727 blinded author ambiguity audit and a 50-cell blinded coefficient-extraction audit were per-
 728 formed after the automated benchmark freeze; neither was used for training, prompt tuning,
 729 model selection, or threshold tuning. The labeling protocols and instructions are released
 730 in the repository.

731 **15. Institutional review board (IRB) approvals or equivalent for research with human**
 732 **subjects**

733 Question: Does the paper describe potential risks incurred by study participants, whether
 734 compensation was provided, etc.?

735 Answer: [\[N/A\]](#)

736 Justification: No human subjects research was conducted.

737 **16. Declaration of LLM usage**

738 Question: Does the paper describe the usage of LLMs in the research process?

739 Answer: [\[Yes\]](#)

740 Justification: LLMs are the subject of evaluation and are used in the benchmark pipeline
 741 for Exp A consensus labeling, field reconstruction, coefficient extraction, and calibration
 742 elicitation. The release includes prompts, model outputs, caches, and audit records needed
 743 to inspect those uses. The authors also used LLM assistance during coding and manuscript
 744 editing, with final responsibility for claims and text retained by the authors.